

4 points to cause a deadlock

Safe state:

- No deadlocks will appear

Avoidance algorithm – for single instance of resource

Multiple resources and multiple instances

- Use banker's algorithm
  - o Each process declares the max # of resources needed to finish
  - o Can force a process to wait (decline giving a resource)
  - o When process gets all resources must return them in a finite amount of time
    - Won't keep all resources forever

Assume # of processes and # of resources (of different types)

Tensor (not matrix)

Max

- nxm matrix
  - o Max # resources for each process

Allocation matrix

- Current allocation between processes and resources
- Define allocation to current process

Need

- Need = max – allocation
- How much left we need to complete the job
- We **compute** this matrix based on the other ones
- In banker algo
  - o Has available, max for each processes, allocation for resources for different system processes

Close relationship between these 3 matrices

Safety Algorithm

- Two different vectors
  - o Work is resources available at current time for all process running on system
  - o Finish – boolean vector
    - False at start
      - No process finished its task
      - If a process finishes task, goes from false to true

Work is available resources on system

- Process completed task, if this conditional is true
  - o Then it goes and sets finish to true

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively.  
Initialize:  

$$\mathbf{Work} = \mathbf{Available}$$

$$\mathbf{Finish}[i] = \text{false for } i = 0, 1, \dots, n-1$$
2. Find an  $i$  such that both:  
 (a)  $\mathbf{Finish}[i] = \text{false}$   
 (b)  $\mathbf{Need}_i \leq \mathbf{Work}$   
 If no such  $i$  exists, go to step 4
3.  $\mathbf{Work} = \mathbf{Work} + \mathbf{Allocation}_i$   
 $\mathbf{Finish}[i] = \text{true}$   
 go to step 2
4. If  $\mathbf{Finish}[i] == \text{true}$  for all  $i$ , then the system is in a safe state

- If not completed
  - o Then we allocate resources to a process

function of sequence of processes in certain order

- Never be in a non-safe state

If we go to stage 4 but finished == false

- We aren't finished, rather we are in a deadlock

Problem

- 5 processes at once
- 5 types of resources
  - o Usb – a – 10 instances (10 possibilities to connect)
  - o Terminal – 5 instances – b
  - o Keyboard or RAM- c (7 instances)

Each process should claim max # resources for each type of resources

- 5 rows (processes), 3 columns (resources)
- We know alloc and max
  - o Compute the need matrix
    - Max – allocation = need
    - Max # resources to complete task
      - P0 needs all these to complete task
        - o Without that , wont complete
        - o Hard constraint on system
          - How do we guess the future?
            - With dynamic allocation, etc we don't know

Work = available vector in beginning

Algo runs before we allocate

- Finds safe sequence of processes to allow it to complete
  - o Safe sequence isn't unique
    - Multiple sequences can be valid

Find = search algorithm

- Limited
- 2<sup>nd</sup> step
  - o Search from first or last?
    - Based on search can find different types of sequences

Algorithm

- Find could be unsuccessful for some search or get a different sequence or order

Prove system is not in a safe state?

- If we find a sequence (allows to run processes in certain order) for the processes = safe
- No sequence exists = unsafe

In reality

- Algorithm (run before we make the decision)
  - o First finds order of processes on system
    - If found safe sequence
      - Then OS will follow the occurrence
    - o As we are sure system won't be in a deadlock

Sometimes one process at a time asks for request

- Now we need to use the request algorithm
  - o Want to be sure that we can allocate for request and not be in a deadlock
    - 2 states

Request > than what we have = deadlock

If we accept request, remain in safe state

- Run the banker algorithm again to find safe state

Only satisfy request if we have enough resources

- If we deny the request, don't need to run the bankers algorithm
- **Necessary condition**
  - o If we accept, are we in a safe state... still need to check the banker algorithm

What changed:

- P1 requests 1 0 2 resources
  - o Can accept
    - If we do, ensure that we still have a safe state

old allocation (OG BA)

can request (3,3,0) by p4, not enough resources

- So not accept

Can we accept (0,2,0) by p0

- Need to check req <= available
  - o If so, need to check
    - Req <= need
      - If true then good



### Resource-Request Algorithm for Process $P_i$

$Request_i$  = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $Request_i \leq Need_i$ , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$Available = Available - Request_i;$

$Allocation_i = Allocation_i + Request_i;$

$Need_i = Need_i - Request_i;$

- o If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- o If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored



### Example: $P_1$ Request (1,0,2)

- o Check that  $Request \leq Available$  (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true

	Allocation	Need	Available
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- o Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- o Can request for (3,3,0) by  $P_4$  be granted?
- o Can request for (0,2,0) by  $P_0$  be granted?

	Allocation	Max	Available
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

New need = old max – new allocation (if we accept new request)

Deadlock detection = new algorithm

- Allow system to enter deadlock
  - o Why?
- If yes, then detect if we are in deadlock or not
  - o If we are = recovery scheme

circular wait = unsafe state

$n^2$  algorithm op

high # of processes on foreground and background

- Graph becomes complex
  - o So use a wait for graph
    - When process waits for another one

Hard to find the graph on left and read it

Lets simplify the concept and detect deadlock via 2<sup>nd</sup> graph

- P4 waits for p1
  - o P1 waits for p2
    - P2 waits for p3
      - Which waits for p4
        - o Do we have a cycle – if yes (deadlock/unsafe)
          - Unsafe = includes both positions
            - As resource can have more than 1 instance
          - P1, p2, p4, p1
          - And a bigger cycle
            - P1,p2,p3,p4,p1
            - So system is in deadlock

Detection algo


3 variables (for each process/type) resource

- Available
  - o Each resource has x instances
- Allocation
- Requests



## Single Instance of Each Resource Type

- Maintain **wait-for** graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph



Resource-Allocation Graph      Corresponding wait-for graph



## Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If **Request**  $[i][j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

- Of all current processes

$N$  = processes

$M$  = resource #

- Each resource can have more than once instance

define 2 more vars

- Set finish value according to allocation
  - Not zero
    - Finish = false
      - Otherwise true
      - Different initialization state

Stage 2

- These conditions need to be agreed on, then go to next state

If doesn't have allocation, doesn't request, so assume it finished

- Each process completes task in certain amount of time

Get to safe state

- Order of sequences
  - 1b = shortcut / optimization

Stage 4

- Meaning is we are in a deadlock
  - Didn't satisfy condition of allocation
  - Unsafe state

goal is to get the same back In work as initial total

4<sup>th</sup> stage, CHECK AND VERIFY THAT FINISH IS TRUE FOR ALL processes

Similar to banker

- A process can ask for additional requests
- Cant allow p2 to run and get 1 byte more
  - Delays everything else
    - Computer crashes because of deadlock
      - Not safe state

Hard to do all checks

- Makes system heavy
- 2 types of OS
  - Unix



## Detection Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively  
Initialize:
  - (a) **Work** = **Available**
  - (b) For  $i = 1, 2, \dots, n$ , if **Allocation** <sub>$i$</sub>   $\neq 0$ , then **Finish**[ $i$ ] = **false**; otherwise, **Finish**[ $i$ ] = **true**
2. Find an index  $i$  such that both:
  - (a) **Finish**[ $i$ ] == **false**
  - (b) **Request** <sub>$i$</sub>   $\leq$  **Work**
 If no such  $i$  exists, go to step 4



## Detection Algorithm (Cont.)

3. **Work** = **Work** + **Allocation** <sub>$i$</sub> ,  
**Finish**[ $i$ ] = **true**  
go to step 2
4. If **Finish**[ $i$ ] == **false**, for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if **Finish**[ $i$ ] == **false**, then  $P_i$  is deadlocked

Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state



## Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)

- Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1		1	0			
$P_4$	0	0	2		0	2			

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in **Finish**[ $i$ ] = **true** for all  $i$

- $P_2$  requests an additional instance of type C

	<u>Request</u>		
	A	B	C
$P_0$	0	0	0
$P_1$	2	0	2
$P_2$	0	0	1
$P_3$		1	0
$P_4$	0	0	2

- State of system?

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes' requests
- Deadlock exists, consisting of processes  $P_1, P_2, P_3$ , and  $P_4$

$P_0$	0, 0, 2	0, 0, 0	KT
$P_1$	5, 2, 6	2, 0, 0	F
$P_2$	7, 2, 6	0, 0, 1	KT
$P_3$	2, 1, 1	0, 1, 0	
$P_4$	0, 0, 2	0, 0, 2	

- Probability is small, don't care
- Windows
  - Can happen
    - Safe mode
    - Sure that every process on system can accomplish task
      - Slow and not easy
        - Can solve computer issues
          - Then return to the regular mode

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ▶ one for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes "caused" the deadlock.

## Recovery from deadlock

- Abort processes
  - Fixes it
- Chose victim
  - Victim usually asks for lots of resources
    - Look at memory allocation
      - If small, don't kill him (doesn't pay off)
  - 103 processes running
    - Pick one that asks for lots of resources
  - But can be important and runs for a long time
    - So hard to say which victim to pick (depends on priority)
      - When all done, select victim process to perform task
        - Runs fast w/o problems



### Recovery from Deadlock: Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
  1. Priority of the process
  2. How long process has computed, and how much longer to completion
  3. Resources the process has used
  4. Resources process needs to complete
  5. How many processes will need to be terminated
  6. Is process interactive or batch?

we define the cost based on task importance

- Based on load or importance
  - Select victim 0 return to safe state (no deadlock)
  - Rollback = restart process we select for that state
    - Don't always select same process as a victim = will be starved and never runs

Select a victim

- Take p2 for example
  - Accept request and acceptance causes deadlock
    - So take p2, p1, or p3
      - P2 seems to be high priority process so maybe not a good idea
    - Maybe p3 or p4
      - We act on the Request part (add 1 more resource)
        - Type C resource
          - They hold one or 2 instances of that type
            - Overcome deadlock problem

Ensure to rollback the process we select



### Recovery from Deadlock: Resource Preemption

- **Selecting a victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor

- Don't take the same person always
- 
- Look at cpu, disk, memory as resources
  - o Don't pick task manager to end as it causes other stuff to go awry (high cost)
    - We don't need edge, lets kill it
      - But rerun it

Same in linux too

A (7 instances), B (2 instances), and C (6 instances)

Name	Status	CPU	Memory	Disk	Network
Service Host: AV/CTP service		0%	1.3 MB	0 MB/s	0 Mbps
Service Host: Background Inte...		0%	7.1 MB	0 MB/s	0 Mbps
Service Host: Capability Acces...		0%	2.7 MB	0 MB/s	0 Mbps
Service Host: Certificate Propa...		0%	1.9 MB	0 MB/s	0 Mbps
Service Host: Clipboard User S...		0%	2.3 MB	0 MB/s	0 Mbps
Service Host: COM+ Event Sys...		0%	2.3 MB	0 MB/s	0 Mbps
Service Host: Connected Devi...		0%	3.7 MB	0 MB/s	0 Mbps
Service Host: Connected Devi...		0%	7.5 MB	0 MB/s	0 Mbps
Service Host: Data Sharing Ser...		0%	3.1 MB	0 MB/s	0 Mbps
Service Host: DCOM Server Pr...		0%	8.7 MB	0 MB/s	0 Mbps
Service Host: Device Associati...		0%	1.1 MB	0 MB/s	0 Mbps
Service Host: DHCP Client		0%	1.8 MB	0 MB/s	0 Mbps
Service Host: Diagnostic Polic...		0%	22.3 MB	0 MB/s	0 Mbps
Service Host: Display Enhance...		0%	1.6 MB	0 MB/s	0 Mbps
Service Host: Display Policy Se...		0%	1.0 MB	0 MB/s	0 Mbps
Service Host: Distributed Link ...		0%	0.8 MB	0 MB/s	0 Mbps